

The Fibonacci sequence

Eric Martin, CSE, UNSW

COMP9021 Principles of Programming

```
[1]: from functools import lru_cache
     from math import sqrt
```

The Fibonacci sequence, say $(F_n)_{n \in \mathbb{N}}$, is defined as $F_0 = 0$, $F_1 = 1$ and for all $n > 1$, $F_n = F_{n-2} + F_{n-1}$; so it is 0, 1, 1, 2, 3, 5, 8, 13, 21, 34...

A generator function is the best option to generate the initial segment of the Fibonacci sequence of a given length, even though it can also be used to generate the member of the Fibonacci sequence of a given index:

```
[2]: def fibonacci_sequence():
     yield 0
     yield 1
     previous, current = 0, 1
     while True:
         previous, current = current, previous + current
         yield current
```

```
[3]: S = fibonacci_sequence()
     list(next(S) for _ in range(19))
```

```
[3]: [0, 1, 1, 2, 3, 5, 8, 13, 21, 34, 55, 89, 144, 233, 377, 610, 987, 1597, 2584]
```

```
[4]: from IPython.display import clear_output

     S = fibonacci_sequence()
     for _ in range(18):
         next(S)
         clear_output()
     next(S)
```

```
[4]: 2584
```

In case only one or a few specific members of the Fibonacci sequence are needed, a simple function is more appropriate:

```
[5]: def iterative_fibonacci(n):
     if n < 2:
         return n
```

```

previous, current = 0, 1
for _ in range(2, n + 1):
    previous, current = current, previous + current
return current

iterative_fibonacci(18)

```

[5]: 2584

A naive recursive implementation is elegant, but too inefficient, as we will see:

```

[6]: def recursive_fibonacci(n):
    if n >= 2:
        return recursive_fibonacci(n - 2) + recursive_fibonacci(n - 1)
    return n

recursive_fibonacci(18)

```

[6]: 2584

Let an integer n greater than 1 be given. Then a call to `recursive_fibonacci(n)` involves:

- for all nonzero $k \leq n$, F_{n-k+1} calls to `recursive_fibonacci(k)`;
- F_{n-1} calls to `recursive_fibonacci(0)`.

In particular, `recursive_fibonacci(n)` calls `recursive_fibonacci(1)` F_n times. Proof is by induction on $k \leq n$:

- `recursive_fibonacci(n)` is called once indeed.
- `recursive_fibonacci(n)` directly calls `recursive_fibonacci(n - 1)` and does not call it indirectly, so calls it once indeed.
- For all $k < n$, `recursive_fibonacci(n - k)` is directly called by `recursive_fibonacci(n - k + 1)` and by `recursive_fibonacci(n - k + 2)`. By inductive hypothesis, the latter two are called directly or indirectly by `recursive_fibonacci(n)` F_k and F_{k-1} times, respectively. Hence `recursive_fibonacci(n - k)` is called by `recursive_fibonacci(n)` F_{k+1} times.
- `recursive_fibonacci(0)` is directly called by `recursive_fibonacci(2)`, hence it is called by `recursive_fibonacci(n)` F_{n-1} times.

Let us illustrate this for $n = 6$ with the following tracing function:

```

[7]: def trace_recursive_fibonacci(n, depth):
    print('    ' * depth, 'Start of function call for n =', n)
    if n >= 2:
        second_previous = trace_recursive_fibonacci(n - 2, depth + 1)
        previous = trace_recursive_fibonacci(n - 1, depth + 1)
        print('    ' * depth, f'End of function call for n = {n}, returning',
              second_previous + previous
              )
    return second_previous + previous
    print('    ' * depth, f'End of function call for n = {n}, returning', n)

```

```
return n
```

```
trace_recursive_fibonacci(6, 0)
```

```
Start of function call for n = 6
  Start of function call for n = 4
    Start of function call for n = 2
      Start of function call for n = 0
      End of function call for n = 0, returning 0
      Start of function call for n = 1
      End of function call for n = 1, returning 1
    End of function call for n = 2, returning 1
    Start of function call for n = 3
      Start of function call for n = 1
      End of function call for n = 1, returning 1
      Start of function call for n = 2
        Start of function call for n = 0
        End of function call for n = 0, returning 0
        Start of function call for n = 1
        End of function call for n = 1, returning 1
      End of function call for n = 2, returning 1
    End of function call for n = 3, returning 2
  End of function call for n = 4, returning 3
  Start of function call for n = 5
    Start of function call for n = 3
      Start of function call for n = 1
      End of function call for n = 1, returning 1
      Start of function call for n = 2
        Start of function call for n = 0
        End of function call for n = 0, returning 0
        Start of function call for n = 1
        End of function call for n = 1, returning 1
      End of function call for n = 2, returning 1
    End of function call for n = 3, returning 2
  End of function call for n = 5, returning 5
  Start of function call for n = 4
    Start of function call for n = 2
      Start of function call for n = 0
      End of function call for n = 0, returning 0
      Start of function call for n = 1
      End of function call for n = 1, returning 1
    End of function call for n = 2, returning 1
    Start of function call for n = 3
      Start of function call for n = 1
      End of function call for n = 1, returning 1
      Start of function call for n = 2
        Start of function call for n = 0
        End of function call for n = 0, returning 0
```

```

        Start of function call for n = 1
    End of function call for n = 1, returning 1
End of function call for n = 2, returning 1
End of function call for n = 3, returning 2
End of function call for n = 4, returning 3
End of function call for n = 5, returning 5
End of function call for n = 6, returning 8

```

[7]: 8

We can still save the recursive design by saving terms of the Fibonacci sequence as they get computed for the first time. As a result of processing the `def` statement below, a dictionary, `fibonacci`, is created and initialised with the values of the first two members of the Fibonacci sequence. Then the function `memoise_fibonacci()` is called, directly as `memoise_fibonacci(18)`, and indirectly as `memoise_fibonacci(18)` executes. For each of those calls, `memoise_fibonacci()` is given one argument only, so the second argument is set to its default value, namely, `fibonacci`, extended with a new key and associated value in case the condition of the `if` statement in the body of `memoise_fibonacci()` evaluates to `True`:

```

[8]: def memoise_fibonacci(n, fibonacci={0: 0, 1: 1}):
    if n not in fibonacci:
        fibonacci[n] = memoise_fibonacci(n - 2) + memoise_fibonacci(n - 1)
    return fibonacci[n]

memoise_fibonacci(18)

```

[8]: 2584

Let us illustrate the mechanism for $n = 6$ with the following tracing function:

```

[9]: def trace_memoise_fibonacci(n, depth, fibonacci={0: 0, 1: 1}):
    print('    ' * depth, 'Start of function call for n =', n)
    print('    ' * (depth + 1), f'fibonacci now is {fibonacci}; ', end = '')
    if n not in fibonacci:
        print('compute value')
        fibonacci[n] = trace_memoise_fibonacci(n - 2, depth + 1)\
            + trace_memoise_fibonacci(n - 1, depth + 1)
    else:
        print('retrieve value')
    print('    ' * depth, f'End of function call for n = {n}, returning',
        fibonacci[n]
    )
    return fibonacci[n]

trace_memoise_fibonacci(6, 0)

```

```

Start of function call for n = 6
    fibonacci now is {0: 0, 1: 1}; compute value

```

```

Start of function call for n = 4
  fibonacci now is {0: 0, 1: 1}; compute value
  Start of function call for n = 2
    fibonacci now is {0: 0, 1: 1}; compute value
    Start of function call for n = 0
      fibonacci now is {0: 0, 1: 1}; retrieve value
    End of function call for n = 0, returning 0
    Start of function call for n = 1
      fibonacci now is {0: 0, 1: 1}; retrieve value
    End of function call for n = 1, returning 1
  End of function call for n = 2, returning 1
  Start of function call for n = 3
    fibonacci now is {0: 0, 1: 1, 2: 1}; compute value
    Start of function call for n = 1
      fibonacci now is {0: 0, 1: 1, 2: 1}; retrieve value
    End of function call for n = 1, returning 1
    Start of function call for n = 2
      fibonacci now is {0: 0, 1: 1, 2: 1}; retrieve value
    End of function call for n = 2, returning 1
  End of function call for n = 3, returning 2
End of function call for n = 4, returning 3
Start of function call for n = 5
  fibonacci now is {0: 0, 1: 1, 2: 1, 3: 2, 4: 3}; compute value
  Start of function call for n = 3
    fibonacci now is {0: 0, 1: 1, 2: 1, 3: 2, 4: 3}; retrieve value
  End of function call for n = 3, returning 2
  Start of function call for n = 4
    fibonacci now is {0: 0, 1: 1, 2: 1, 3: 2, 4: 3}; retrieve value
  End of function call for n = 4, returning 3
End of function call for n = 5, returning 5
End of function call for n = 6, returning 8

```

[9]: 8

`memoise_fibonacci()` illustrates the fact that when a function argument has a default value, that default value is not created at every function call, but at the time when Python processes the function's `def` statement. This makes no difference for default values of a type such as `int`:

```

[10]: def f(x=0):
      x += 1
      return x

# Create the argument 0 before calling f(), let x denote it, from the
# value denoted by x and 1 create 1, let x denote it.
f(0)
f(1)
f(2)
# Let x denote the 0 created when def was processed, from the value

```

```
# denoted by x and 1 create 1, let x denote it.  
f()  
f()  
f()
```

[10]: 1

[10]: 2

[10]: 3

[10]: 1

[10]: 1

[10]: 1

But it makes a difference for default values of a type such as `list`:

```
[11]: def g(x=[0]):  
      x += [1]  
      return x  
  
# Create the argument [0] before calling g(), let x denote it, then  
# extend it to [0, 1], let x denote the modified list.  
g([0])  
g([1])  
g([2])  
# Let x denote the list L created when def was processed, then and now  
# equal to [0], then extend it to [0, 1], let x denote the modified L.  
g()  
# Let x denote the list L created when def was processed, now equal to  
# [0, 1], then extend it to [0, 1, 1], let x denote the modified L.  
g()  
g()
```

[11]: [0, 1]

[11]: [1, 1]

[11]: [2, 1]

[11]: [0, 1]

[11]: [0, 1, 1]

[11]: [0, 1, 1, 1]

What was good for `memoise_fibonacci()` might not be the intended behaviour for other functions, in other contexts: in case a function F is called without an argument for a parameter p that in F 's definition, receives a default value v , one might want p to always be assigned that default value, not the value currently denoted by p and possibly modified from the original value of v following previous calls to F . One should then opt for a different design:

```
[12]: def h(x=None):
      if x is None:
          x = [0]
      x += [1]
      return x

      # Create the argument [0] before calling h(), let x denote it, then
      # extend it to [0, 1], let x denote the modified list.
h([0])
h([1])
h([2])
      # Let x denote None, then create [0], let x denote it, then extend it to
      # [0, 1], let x denote the modified list.
h()
h()
h()
```

```
[12]: [0, 1]
```

```
[12]: [1, 1]
```

```
[12]: [2, 1]
```

```
[12]: [0, 1]
```

```
[12]: [0, 1]
```

```
[12]: [0, 1]
```

The `lru_cache()` (“lru” is for *Least Recently Used*) function from the `functools` module returns a function that can be used as a **decorator** and applied to a function F to yield a memoised version of F . By default, the `maxsize` argument of `lru_cache()` is set to 128, to record up to the last 128 computed values of the function, as witnessed by the `cache_info()` attribute of the memoised version of f :

```
[13]: @lru_cache()
      def lru_fibonacci(n):
          if n < 2:
              return n
          return lru_fibonacci(n - 1) + lru_fibonacci(n - 2)

      lru_fibonacci.cache_info()
```

```
[13]: CacheInfo(hits=0, misses=0, maxsize=128, currsize=0)
```

Suppose that `lru_fibonacci()` is called for the first time with 2 as argument. Since `lru_fibonacci(2)` has not been computed yet, `lru_fibonacci(1)` and `lru_fibonacci(0)` are called, which both have not been computed yet either: a total of 3 values fail to be retrieved (3 misses). The last two values are computed and recorded, then the former value is computed and recorded, and the cache eventually stores those 3 values:

```
[14]: lru_fibonacci(2)
lru_fibonacci.cache_info()
```

```
[14]: 1
```

```
[14]: CacheInfo(hits=0, misses=3, maxsize=128, currsize=3)
```

Calling `lru_fibonacci(2)` again, the value is found in the cache (1 hit):

```
[15]: lru_fibonacci(2)
lru_fibonacci.cache_info()
```

```
[15]: 1
```

```
[15]: CacheInfo(hits=1, misses=3, maxsize=128, currsize=3)
```

When calling `lru_fibonacci(3)`, the value fails to be found in the cache (1 more miss), so `lru_fibonacci(2)` and `lru_fibonacci(1)` are called and retrieved from the cache (2 more hits), and the computed value of `lru_fibonacci(3)` is added to the cache:

```
[16]: lru_fibonacci(3)
lru_fibonacci.cache_info()
```

```
[16]: 2
```

```
[16]: CacheInfo(hits=3, misses=4, maxsize=128, currsize=4)
```

The cache can be cleared with the `cache_clear()` attribute of the memoised version of the function. Then calling `lru_fibonacci(3)` necessitates to call `lru_fibonacci(2)` and `lru_fibonacci(1)`, calling `lru_fibonacci(2)` necessitates to call `lru_fibonacci(1)` and `lru_fibonacci(0)`, for a total of 4 misses that are computed and all stored in the cache:

```
[17]: lru_fibonacci.cache_clear()
lru_fibonacci(3)
lru_fibonacci.cache_info()
```

```
[17]: 2
```

```
[17]: CacheInfo(hits=1, misses=4, maxsize=128, currsize=4)
```


Clearing the cache again, calling `lru_fibonacci(128)` necessitates to call for the first time `lru_fibonacci(128)`, ..., `lru_fibonacci(0)` (129 misses). When calling `lru_fibonacci(2)` for the first time, `lru_fibonacci(1)` could be called before `lru_fibonacci(0)` or the other way around. Execution of the following cell reveals that `lru_fibonacci(0)` is called first; its value leaves the cache after the values of `lru_fibonacci(1)`, ..., `lru_fibonacci(128)` have then been computed and recorded. When `lru_fibonacci(3)` is computed, `lru_fibonacci(1)` is retrieved (whether `lru_fibonacci(1)` or `lru_fibonacci(2)` is computed first), ..., when `lru_fibonacci(128)` is computed, `lru_fibonacci(126)` is retrieved (whether `lru_fibonacci(126)` or `lru_fibonacci(127)` is computed first), for a total of 126 hits:

```
[18]: lru_fibonacci.cache_clear()
lru_fibonacci(128)
lru_fibonacci.cache_info()
lru_fibonacci(1)
lru_fibonacci.cache_info()
lru_fibonacci(0)
lru_fibonacci.cache_info()
```

```
[18]: 251728825683549488150424261
```

```
[18]: CacheInfo(hits=126, misses=129, maxsize=128, currsize=128)
```

```
[18]: 1
```

```
[18]: CacheInfo(hits=127, misses=129, maxsize=128, currsize=128)
```

```
[18]: 0
```

```
[18]: CacheInfo(hits=127, misses=130, maxsize=128, currsize=128)
```

The capacity of the cache can be left unbounded by setting the value of the `maxsize` argument of `lru_cache()` to `None`:

```
[19]: @lru_cache(None)
def unbounded_lru_fibonacci(n):
    if n < 2:
        return n
    return unbounded_lru_fibonacci(n - 1) + unbounded_lru_fibonacci(n - 2)
```

```
[20]: unbounded_lru_fibonacci(150)
unbounded_lru_fibonacci.cache_info()
```

```
[20]: 9969216677189303386214405760200
```

```
[20]: CacheInfo(hits=148, misses=151, maxsize=None, currsize=151)
```

The argument `maxsize` of `lru_cache()` can also be set to any integer value. Let us set it to 4 and first call `bounded_lru_fibonacci(8)`. Then `bounded_lru_fibonacci(8)`, `bounded_lru_fibonacci(7)`, `bounded_lru_fibonacci(6)` and `bounded_lru_fibonacci(5)` are

last called and recorded. If `bounded_lru_fibonacci(5)` is then called, its value is retrieved (1 more hit). And if `bounded_lru_fibonacci(4)` is thereafter called, `bounded_lru_fibonacci(4)`, ..., `bounded_lru_fibonacci(0)` have to be recomputed (5 more misses), with `bounded_lru_fibonacci(2)` and `bounded_lru_fibonacci(1)` being retrieved in the process (2 more hits):

```
[21]: @lru_cache(4)
def bounded_lru_fibonacci(n):
    if n < 2:
        return n
    return bounded_lru_fibonacci(n - 1) + bounded_lru_fibonacci(n - 2)
```

```
[22]: bounded_lru_fibonacci(8)
bounded_lru_fibonacci.cache_info()
bounded_lru_fibonacci(5)
bounded_lru_fibonacci.cache_info()
bounded_lru_fibonacci(4)
bounded_lru_fibonacci.cache_info()
```

[22]: 21

[22]: CacheInfo(hits=6, misses=9, maxsize=4, currsz=4)

[22]: 5

[22]: CacheInfo(hits=7, misses=9, maxsize=4, currsz=4)

[22]: 3

[22]: CacheInfo(hits=9, misses=14, maxsize=4, currsz=4)

Set $\varphi = \frac{1+\sqrt{5}}{2}$ and $\psi = \frac{1-\sqrt{5}}{2}$. For all $n \in \mathbb{N}$, $(\frac{1+\sqrt{5}}{2})^{n+2} = (\frac{1+\sqrt{5}}{2})^n \frac{1+2\sqrt{5}+5}{4} = (\frac{1+\sqrt{5}}{2})^n + (\frac{1+\sqrt{5}}{2})^n \frac{1+\sqrt{5}}{2}$, hence for $x = \varphi$ or $x = \psi$, $x^{n+2} = x^n + x^{n+1}$. Let integers a and b be given, and for all $n \in \mathbb{N}$, let s_n denote $a\varphi^n + b\psi^n$. It follows from the previous equalities (one for $x = \varphi$, one for $x = \psi$) that for all $n \in \mathbb{N}$, $s_{n+2} = s_n + s_{n+1}$. So $(s_n)_{n \in \mathbb{N}} = (F_n)_{n \in \mathbb{N}}$ iff $s_0 = 0$ and $s_1 = 1$, which is equivalent to the two equalities $a + b = 0$ and $a\varphi + b\psi = 0$, which is equivalent to $a = \frac{1}{\varphi - \psi}$ and $b = -a$, so $a = \frac{1}{\sqrt{5}}$ and $b = -\frac{1}{\sqrt{5}}$. Hence for all $n \in \mathbb{N}$,

$$F_n = \frac{1}{\sqrt{5}} \left[\left(\frac{1+\sqrt{5}}{2} \right)^n - \left(\frac{1-\sqrt{5}}{2} \right)^n \right]$$

This is the *closed-form expression* of the Fibonacci numbers. Note that $|\frac{1}{\sqrt{5}} \frac{1-\sqrt{5}}{2}| < \frac{1}{2}$, hence F_n can be computed as $\frac{1}{\sqrt{5}} (\frac{1+\sqrt{5}}{2})^n$ rounded to the closest integer, resulting in a simpler calculation:

```
[23]: def closed_form_fibonacci(n):
    sqrt_5 = sqrt(5)
    return round(1 / sqrt_5 * ((1 + sqrt_5) / 2) ** n)
```

But of course, due to the limited precision of floating point computation, it does not need a large input for `closed_form_fibonacci()` to fail and produce the correct result:

```
[24]: n = 0
correct_sequence = fibonacci_sequence()
while True:
    correct = next(correct_sequence)
    maybe_incorrect = closed_form_fibonacci(n)
    if maybe_incorrect != correct:
        print(f'{n}th term of the sequence is {correct}, '
              f'incorrectly computed as {maybe_incorrect}.')
    )
    break
    n += 1
```

71th term of the sequence is 308061521170129, incorrectly computed as 308061521170130.